

on specific models: The R package `pre`¹⁴ implements **RuleFit** and H-statistic. The R package `gbm`¹⁵ implements gradient boosted models and H-statistic.

7.3.7 Alternatives

The H-statistic is not the only way to measure interactions:

Variable Interaction Networks (VIN) by Hooker (2004)¹⁶ is an approach that decomposes the prediction function into main effects and feature interactions. The interactions between features are then visualized as a network. Unfortunately no software is available yet.

Partial dependence based feature interaction by Greenwell et al. (2018)¹⁷ measures the interaction between two features. This approach measures the feature importance (defined as the variance of the partial dependence function) of one feature conditional on different, fixed points of the other feature. If the variance is high, then the features interact with each other, if it is zero, they do not interact. The corresponding R package `vip` is available on Github¹⁸. The package also covers partial dependence plots and feature importance.

7.4 Functional Decomposition

A supervised machine learning model can be viewed as a function that takes a high-dimensional feature vector as input and produces a prediction or classification score as output. Functional decomposition is an interpretation technique that deconstructs the high-dimensional function and expresses it as a sum of individual feature effects and interaction effects that can be visualized. In addition, functional decomposition is a fundamental principle underlying many interpretation techniques – it helps you better understand other interpretation methods.

Let us jump right in and look at a particular function. This function takes two features as input and produces a one-dimensional output:

$$y = \hat{f}(x_1, x_2) = 2 + e^{x_1} - x_2 + x_1 \cdot x_2$$

¹⁴<https://cran.r-project.org/web/packages/pre/index.html>

¹⁵<https://github.com/gbm-developers/gbm3>

¹⁶Hooker, Giles. “Discovering additive structure in black box functions.” Proceedings of the tenth ACM SIGKDD international conference on Knowledge discovery and data mining. (2004).

¹⁷Greenwell, Brandon M., Bradley C. Boehmke, and Andrew J. McCarthy. “A simple and effective model-based variable importance measure.” arXiv preprint arXiv:1805.04755 (2018).

¹⁸<https://github.com/koalaverse/vip>

2 features 的 PDP 可以反映
features 之间的关系。

supervised model can
be viewed as a function.

deconstructs
high-dimensional function
and express it as a sum
of individual feature effects

Think of the function as a machine learning model. We can visualize the function with a 3D plot or a heatmap with contour lines:




Figure 7.22: Prediction surface of a function with two features X_1 and X_2 .

The function takes large values when X_1 is large and X_2 is small, and it takes small values for large X_2 and small X_1 . The prediction function is not simply an additive effect between the two features, but an interaction between the two. The presence of an interaction can be seen in the figure – the effect of changing values for feature X_1 depends on the value that feature X_2 has.

Our job now is to decompose this function into main effects of features X_1 and X_2 and an interaction term. For a two-dimensional function \hat{f} that depends on only two input features: $\hat{f}(x_1, x_2)$, we want each component to represent a main effect (\hat{f}_1 and \hat{f}_2), interaction ($\hat{f}_{1,2}$) or **intercept** (\hat{f}_0):

$$\hat{f}(x_1, x_2) = \hat{f}_0 + \hat{f}_1(x_1) + \hat{f}_2(x_2) + \hat{f}_{1,2}(x_1, x_2)$$

The main effects indicate how each feature affects the prediction, independent of the

main effect : \hat{f}_1 and \hat{f}_2

interaction = $\hat{f}_{1,2}$

intercept : \hat{f}_0

截距

values the other feature. The interaction effect indicates the joint effect of the features. The intercept simply tells us what the prediction is when all feature effects are set to zero. Note that the components themselves are functions (except for the intercept) with different input dimensionality.

I will just give you the components now and explain where they come from later. The intercept is $\hat{f}_0 \sim 3.18$. Since the other components are functions, we can visualize them:




Figure 7.23: Decomposition of a function.

Do you think the components make sense given the above true formula, ignoring that the intercept value seems a bit random? The x_1 feature shows an **exponential** main effect, and x_2 shows a negative linear effect. The interaction term looks a bit like a Pringles chip. In less crunchy and more mathematical terms, it is a **hyperbolic paraboloid**, as we would expect for $x_1 \cdot x_2$. Spoiler alert: the decomposition is based on accumulated local effect plots, which we will discuss later in the chapter.

exponential 指数级
hyperbolic 双曲线的
paraboloid 抛物面

7.4.1 How not to Compute the Components I

But why all the excitement? A glance at the formula already gives us the answer to the decomposition, so no need for fancy methods, right? For feature x_1 , we can take all the **summands** that contain only x_1 as the component for that feature. That would be $\hat{f}_1(x_1) = e^{x_1}$ and $\hat{f}_2(x_2) = -x_2$ for feature x_2 . The interaction is then $\hat{f}_{12}(x_1, x_2) = x_1 \cdot x_2$. While this is the correct answer for this example (up to constants), there are two problems with this approach: Problem 1): While the example started with the formula, the reality is that **almost no machine learning model can be described with such a neat formula**. Problem 2) is much more **intricate** and concerns **what an interaction is**. Imagine a simple function $\hat{f}(x_1, x_2) = x_1 \cdot x_2$, where both features take values larger than zero and are independent of each other. Using our look-at-the-formula **tactic**, we would conclude that there is an interaction between features x_1 and x_2 , but not individual feature effects. But can we really say that feature x_1 has no individual effect on the prediction function? Regardless of what value the other feature x_2 takes on, the prediction increases as we increase x_1 . For example, for $x_2 = 1$, the effect of x_1 is $\hat{f}(x_1, 1) = x_1$, and when $x_2 = 10$ the effect is $\hat{f}(x_1, 10) = 10 \cdot x_1$. Thus, it is clear that feature x_1 has a positive effect on the prediction, independent of x_2 , and is not zero.

To solve problem 1) of lack of access to a neat formula, we need a method that uses only the prediction function or classification score. To solve problem 2) of lack of definition, we need some **axioms** that tell us what the components should look like and how they relate to each other. But first, we should define more precisely what functional decomposition is.

7.4.2 Functional Decomposition

A prediction function takes p features as input, $\hat{f} : \mathbb{R}^p \mapsto \mathbb{R}$ and produces an output. This can be a regression function, but it can also be the classification probability for a given class or the score for a given cluster (unsupervised machine learning). Fully decomposed, we can represent the prediction function as the sum of functional components:

$$\begin{aligned} \hat{f}(x) = & \hat{f}_0 + \hat{f}_1(x_1) + \dots + \hat{f}_p(x_p) \\ & + \hat{f}_{1,2}(x_1, x_2) + \dots + \hat{f}_{1,p}(x_1, x_p) + \dots + \hat{f}_{p-1,p}(x_{p-1}, x_p) \\ & + \dots \\ & + \hat{f}_{1,\dots,p}(x_1, \dots, x_p) \end{aligned}$$

We can make the decomposition formula a bit nicer by indexing all possible subsets of feature combinations: $S \subseteq \{1, \dots, p\}$. This set contains the intercept ($S = \emptyset$), main

excitement 刺激因素

summand 被加数

neat 整齐的;简洁的

intricate 难以理解的

tactic 策略;

Problem 1:

no neat formula for ml model

Problem 2:

what interaction is

axiom 公理

effects ($|S| = 1$), and all interactions ($|S| \geq 1$). With this subset defined, we can write the decomposition as follows:

$$\hat{f}(x) = \sum_{S \subseteq \{1, \dots, p\}} \hat{f}_S(x_S)$$

In the formula, x_S is the vector of features in the index set S . And each subset S represents a functional component, for example a main effect if S contains only one feature, or an interaction if $|S| > 1$.

How many components are in the above formula? The answer boils down to how many possible subsets S of the features $1, \dots, p$ we can form. And these are $\sum_{i=0}^p \binom{p}{i} = 2^p$ possible subsets! For example, if a function uses 10 features, we can decompose the function into 1042 components: 1 intercept, 10 main effects, 90 2-way interaction terms, 720 3-way interaction terms, ... And with each additional feature, the number of components doubles. Clearly, for most functions, it is not feasible to compute all components. Another reason NOT to compute all components is that components with $|S| > 2$ are difficult to visualize and interpret.

7.4.3 How not to Compute the Components II

So far I have avoided talking about how the components are defined and computed. The only constraints we have implicitly talked about were the number and dimensionality of the components, and that the sum of components should yield the original function. But without further constraints on what the components should be, they are not unique. This means we could shift effects between main effects and interactions, or lower-order interactions (few features) and higher-order interactions (more features). In the example at the beginning of the chapter we could set both main effects to zero, and add their effects to the interaction effect.

Here is an even more extreme example that illustrates the need for constraints on components. Suppose you have a 3-dimensional function. It does not really matter what this function looks like, but the following decomposition would **always** work: \hat{f}_0 is 0.12. $\hat{f}_1(x_1) = 2 \cdot x_1 + \text{number of shoes you own}$. $\hat{f}_2, \hat{f}_3, \hat{f}_{1,2}, \hat{f}_{2,3}, \hat{f}_{1,3}$ are all zero. And to make this trick work, I define $\hat{f}_{1,2,3}(x_1, x_2, x_3) = \hat{f}(x) - \sum_{S \subset \{1, \dots, p\}} \hat{f}_S(x_S)$. So the interaction term containing all features simply sucks up all the remaining effects, which by definition always works, in the sense that the sum of all components gives us the original prediction function. This decomposition would not be very meaningful and quite misleading if you were to present this as the interpretation of your model.

boils down 归结为

implicitly 暗中地

dimensionality 维度

yield 产生; 屈服; 让步

formula 的顺序不影响解释

The ambiguity can be avoided by specifying further constraints or specific methods for computing the components. In this chapter, we will discuss three methods that approach the functional decomposition in different ways:

- (Generalized) Functional ANOVA
- **Accumulated Local Effects**
- Statistical regression models

三种方法做 functional decomposition

7.4.4 Functional ANOVA

Functional ANOVA was proposed by Hooker (2004)¹⁹. A requirement for this approach is that the model prediction function \hat{f} is square integrable. As with any functional decomposition, the functional ANOVA decomposes the function into components:

$$\hat{f}(x) = \sum_{S \subseteq \{1, \dots, p\}} \hat{f}_S(x_S)$$

Hooker (2004) defines each component with the following formula:

$$\hat{f}_S(x) = \int_{X_{-S}} \left(\hat{f}(x) - \sum_{V \subset S} \hat{f}_V(x) \right) dX_{-S}$$

Okay, let us take this thing apart. We can rewrite the component as:

$$\hat{f}_S(x) = \int_{X_{-S}} (\hat{f}(x)) dX_{-S} - \int_{X_{-S}} \left(\sum_{V \subset S} \hat{f}_V(x) \right) dX_{-S}$$

integral 积分

On the left side is the **integral** over the prediction function with respect to the **features excluded from the set S , denoted with $-S$** . For example, if we compute the 2-way interaction component for features 2 and 3, we would integrate over features 1, 4, 5, ... The integral can also be viewed as the expected value of the prediction function with respect to X_{-S} , assuming that all features follow a uniform distribution from their minimum to their maximum. From this interval, we subtract all components with subsets of S . This subtraction removes the effect of all lower-order effects and centers the effect. For $S = \{1, 2\}$, we subtract the main effects of both features \hat{f}_1 and \hat{f}_2 , as well as the intercept \hat{f}_0 . The **occurrence** of these lower-order effects makes the formula recursive: We have to go through the hierarchy of subsets to the intercept and compute all these components.

occurrence 出现; 事件(已发生)

¹⁹Hooker, Giles. "Discovering additive structure in black box functions." Proceedings of the tenth ACM SIGKDD international conference on Knowledge discovery and data mining. (2004).

For the intercept component \hat{f}_0 , the subset is the empty set $S = \{\emptyset\}$ and therefore $-S$ contains all features:

$$\hat{f}_0(x) = \int_X \hat{f}(x) dX$$

This is simply the prediction function integrated over all features. The intercept can also be interpreted as the expectation of the prediction function when we assume that all features are uniformly distributed. Now that we know \hat{f}_0 , we can compute \hat{f}_1 (and equivalently \hat{f}_2):

$$\hat{f}_1(x) = \int_{X_{-1}} (\hat{f}(x) - \hat{f}_0) dX_{-S}$$

To finish the calculation for the component $\hat{f}_{1,2}$, we can put everything together:

$$\begin{aligned} \hat{f}_{1,2}(x) &= \int_{X_{3,4}} (\hat{f}(x) - (\hat{f}_0(x) + \hat{f}_1(x) - \hat{f}_0 + \hat{f}_2(x) - \hat{f}_0)) dX_{3,4} \\ &= \int_{X_{3,4}} (\hat{f}(x) - \hat{f}_1(x) - \hat{f}_2(x) + \hat{f}_0) dX_{3,4} \end{aligned}$$

This example shows how each higher-order effect is defined by integrating over all other features, but also by removing all the lower-order effects that are subsets of the feature set we are interested in.

Hooker (2004) has shown that this definition of functional components satisfies these desirable axioms:

- Zero Means: $\int \hat{f}_S(x_S) dX_S = 0$ for each $S \neq \emptyset$.
- Orthogonality: $\int \hat{f}_S(x_S) \hat{f}_V(x_V) dX = 0$ for $S \neq V$
- Variance Decomposition: Let $\sigma_{\hat{f}}^2 = \int \hat{f}(x)^2 dX$, then $\sigma^2(\hat{f}) = \sum_{S \subseteq \{1, \dots, p\}} \sigma_S^2(\hat{f}_S)$

The zero means axiom implies that all effects or interactions are centered around zero. As a consequence, the interpretation at a position x is relative to the centered prediction and not the absolute prediction.

The orthogonality axiom implies that components do not share information. For example, the first-order effect of feature X_1 and the interaction term of X_1 and X_2 are not correlated. Because of orthogonality, all components are “pure” in the sense that they do not mix effects. It makes a lot of sense that the component for, say, feature X_4 should be independent of the interaction term between features X_1 and X_2 . The more interesting

dX_3, X_4 ?

Orthogonality 正交性.

说明任何 component 都是绝对独立的

first order effect - 一阶效应

$\hat{f}(x_1)$ 和 $\hat{f}(x_1, x_2)$

没有任何关联

consequence arises for orthogonality of hierarchical components, where one component contains features of another, for example the interaction between X_1 and X_2 , and the main effect of feature X_1 . In contrast, a two-dimensional partial dependence plot for X_1 and X_2 would contain four effects: the intercept, the two main effects of X_1 and X_2 and the interaction between them. The functional ANOVA component for $\hat{f}_{1,2}(x_1, x_2)$ contains only the pure interaction.

Variance decomposition allows us to divide the variance of the function \hat{f} among the components, and guarantees that it adds up the total variance of the function in the end. The variance decomposition property can also explain to us why the method is called “functional ANOVA”. In statistics, ANOVA stands for ANalysis Of VAriance. ANOVA refers to a collection of methods that analyze differences in the mean of a target variable. ANOVA works by dividing the variance and attributing it to the variables. Functional ANOVA can therefore be seen as an extension of this concept to any function.

Problems arise with the functional ANOVA when features are correlated. As a solution, the generalized functional ANOVA has been proposed.

7.4.5 Generalized Functional ANOVA for Dependent Features

Similar to most interpretation techniques based on sampling data (such as the PDP), the functional ANOVA can produce misleading results when features are correlated. If we integrate over the uniform distribution, when in reality features are dependent, we create a new dataset that deviates from the joint distribution and extrapolates to unlikely combinations of feature values.

Hooker (2007)²⁰ proposed the generalized functional ANOVA, a decomposition that works for dependent features. It is a generalization of the functional ANOVA we encountered earlier, which means that the functional ANOVA is a special case of the generalized functional ANOVA. The components are defined as projections of f onto the space of additive functions:

$$\hat{f}_S(x_S) = \operatorname{argmin}_{g_S \in L^2(\mathbb{R}^S)_{S \in P}} \int \left(\hat{f}(x) - \sum_{S \subset P} g_S(x_S) \right)^2 w(x) dx.$$

Instead of orthogonality, the components satisfy a hierarchical orthogonality condition:

$$\forall \hat{f}_S(x_S) | S \subset U : \int \hat{f}_S(x_S) \hat{f}_U(x_U) w(x) dx = 0$$

²⁰Hooker, Giles. “Generalized functional anova diagnostics for high-dimensional functions of dependent variables.” Journal of Computational and Graphical Statistics 16.3 (2007): 709-732.

当 features 关联时, ANOVA 还是有问题的

deviate { v. 偏离; 使偏离
n. 反常的

extrapolate v. 推断

???

Hierarchical orthogonality is different from orthogonality. For two feature sets S and U , neither of which is the subset of the other (e.g. $S = \{1, 2\}$ and $U = \{2, 3\}$), the components \hat{f}_S and \hat{f}_U need not be orthogonal for the decomposition to be hierarchically orthogonal. But all components for all subsets of S must be orthogonal to \hat{f}_S . As a result, the interpretation differs in relevant ways: Similar to the M-Plot in the [ALE chapter](#), generalized functional ANOVA components can **entangle** the (marginal) effects of correlated features. Whether the components entangle the marginal effects also depends on the choice of weight function $w(x)$. If we choose w to be the uniform measure on the unit cube, we obtain the functional ANOVA from the section above. A natural choice for w is the joint probability distribution function. However, the joint distribution is usually unknown, and difficult to estimate. A trick can be to start with the uniform measure on the unit cube, and cut out areas without data.

entangle 便纠缠; 卷入

The estimation is done on a grid of points in the feature space and is stated as a minimization problem that can be solved using regression techniques. **However, the components cannot be computed individually, nor hierarchically, but a complex system of equations involving other components has to be solved. The computation is therefore quite complex and computationally intensive.**

巨方程

7.4.6 Accumulated Local Effect Plots

ALE plots (Apley and Zhu 2020²¹) also provide a functional decomposition, meaning that adding all ALE plots from intercept, 1D ALE plots, 2D ALE plots and so on, yields the prediction function. ALE differs from the (generalized) functional ANOVA, as the components are not orthogonal but, as the authors call it, **pseudo-orthogonal**. To understand pseudo-orthogonality, we have to define the operator H_S , which takes a function \hat{f} and maps it to its ALE plot for feature subset S . For example, the operator $H_{1,2}$ takes as input a machine learning model and produces the 2D ALE plot for features 1 and 2: $H_{1,2}(\hat{f}) = \hat{f}_{ALE,12}$. If we apply the same operator twice, we get the same ALE plot. After applying the operator $H_{1,2}$ to f once, we get the 2D ALE plot $\hat{f}_{ALE,12}$. Then we apply the operator again, not to f but to $\hat{f}_{ALE,12}$. This is possible because the 2D ALE component is itself a function. The result is again $\hat{f}_{ALE,12}$, meaning we can apply the same operator several times and always get the same ALE plot. This is the first part of pseudo-orthogonality. But what is the result if we apply two different operators for different feature sets? For example, $H_{1,2}$ and H_1 , or $H_{1,2}$ and $H_{3,4,5}$? The answer is zero. If we first apply the ALE operator H_S to a function and then apply H_U to the result (with $S \neq U$), then the result is zero. In other words, the ALE plot of an ALE plot is zero, unless you apply the same ALE plot twice. Or in other words, the ALE plot for the

pseudo-orthogonal 伪正交

²¹Apley, Daniel W., and Jingyu Zhu. "Visualizing the effects of predictor variables in black box supervised learning models." *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 82.4 (2020): 1059-1086.

feature set S does not contain any other ALE plots in it. Or in mathematical terms, the ALE operator maps functions to orthogonal subspaces of an inner product space.

As Apley and Zhu (2020) note, pseudo-orthogonality may be more desirable than hierarchical orthogonality, because it does not entangle marginal effects of the features. Furthermore, ALE does not require estimation of the joint distribution; the components can be estimated in a hierarchical manner, which means that calculating the 2D ALE for features 1 and 2 requires only the calculations of individual ALE components of 1 and 2 and the intercept term in addition.

2D ALE_{1,2} 只计算单独的 ALE₁ 和 ALE₂ 以及截距的和

7.4.7 Statistical Regression Models

This approach ties in with **interpretable models**, in particular **generalized additive models**. Instead of decomposing a complex function, we can build constraints into the modeling process so that we can easily read out the individual components. While decomposition can be handled in a top-down manner, where we start with a high-dimensional function and decompose it, generalized additive models provide a bottom-up approach, where we build the model from simple components. Both approaches have in common that their goal is to provide individual and interpretable components. In statistical models, we restrict the number of components so that not all 2^p components have to be fitted. The simplest version is linear regression:

$$\hat{f}(x) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$$

The formula looks very similar to the functional decomposition, but with two major modifications. Modification 1: All interaction effects are excluded and we keep only intercept and main effects. Modification 2: The main effects may only be linear in the features: $\hat{f}_j(x_j) = \beta_j x_j$. Viewing the linear regression model through the lens of functional decomposition, we see that the model itself represents a functional decomposition of the true function that maps from features to target, but under strong assumptions that the effects are linear effects and there are no interactions.

The generalized additive model relaxes the second assumption by allowing more flexible functions \hat{f}_j through the use of splines. Interactions can also be added, but this process is rather manual. Approaches such as GA2M attempt to add 2-way interactions automatically to a GAM.²²

Thinking of a linear regression model or a GAM as functional decomposition can also lead to confusion. If you apply the decomposition approaches from earlier in the chapter

²²Caruana, Rich, et al. "Intelligible models for healthcare: Predicting pneumonia risk and hospital 30-day readmission." Proceedings of the 21th ACM SIGKDD international conference on knowledge discovery and data mining. (2015).

(generalized functional ANOVA and accumulated local effects), you may get components that are different from the components read directly from the GAM. This can happen when interaction effects of correlated features are modeled in the GAM. The discrepancy occurs because other functional decomposition approaches split effects differently between interactions and main effects.

So when should you use GAMs instead of a complex model + decomposition? You should stick to GAMs when most interactions are zero, especially when there are no interactions with three or more features. If we know that the maximum number of features involved in interactions is two ($|S| \leq 2$), then we can use approaches like MARS or GA2M. Ultimately, model performance on test data may indicate whether a GAM is sufficient or a more complex model performs much better.

7.4.8 Bonus: Partial Dependence Plot

Does the **Partial Dependence Plot** also provide a functional decomposition? Short answer: No. Longer answer: The partial dependence plot for a feature set S always contains all effects of the hierarchy – the PDP for $\{1, 2\}$ contains not only the interaction, but also the individual feature effects. As a consequence, adding all PDPs for all subsets does not yield the original function, and thus is not a valid decomposition. But could we adjust the PDP, perhaps by removing all lower effects? Yes, we could, but we would get something similar to the functional ANOVA. However, instead of integrating over a uniform distribution, the PDP integrates over the marginal distribution of X_{-S} , which is estimated using Monte Carlo sampling.

7.4.9 Advantages

I consider functional decomposition to be a **core concept of machine learning interpretability**.

Functional decomposition gives us a **theoretical justification** for decomposing high-dimensional and complex machine learning models into individual effects and interactions – a necessary step that allows us to interpret individual effects. Functional decomposition is the core idea for techniques such as statistical regression models, ALE, (generalized) functional ANOVA, PDP, the H-statistic, and ICE curves.

Functional decomposition also provides a **better understanding of other methods**. For example, **permutation feature importance** breaks the association between a feature and the target. Viewed through the functional decomposition lens, we can see that the permutation “destroys” the effect of all components in which the feature was involved. This affects the main effect of the feature, but also all interactions with other features. As another example, Shapley values decompose a prediction into additive effects of the

PDP也许能做到 function
但是这个过程会 decomposition
和 ANOVA 相似。

优点:

- ① 它是 XAI 的核心。
- ② 是其它技术的理论核心。
- ③ 有助于理解其它方法。

individual feature. But the functional decomposition tells us that there should also be interaction effects in the decomposition, so where are they? Shapley values provide a fair attribution of effects to the individual features, meaning that all interactions are also fairly attributed to the features and therefore divided up among the Shapley values.

When considering functional decomposition as a tool, the use of **ALE plots** offers many **advantages**. ALE plots provide a functional decomposition that is fast to compute, has software implementations (see the ALE chapter), and desirable pseudo-orthogonality properties.

7.4.10 Disadvantages

The concept of functional decomposition quickly reaches its **limits for high-dimensional components** beyond interactions between two features. Not only does this **exponential explosion** in the number of features limit **practicability**, since we cannot easily visualize higher-order interactions, but computational time is insane if we were to compute all interactions.

Each method of functional decomposition method has their **individual disadvantages**. The bottom-up approach – constructing regression models – is a quite manual process and imposes many constraints on the model that can affect predictive performance. Functional ANOVA requires independent features. Generalized functional ANOVA is very difficult to estimate. Accumulated local effect plots do not provide a variance decomposition.

The functional decomposition approach is **more appropriate for analyzing tabular data than text or images**.

实际上的 feature 在什么地方
Shapley 值说明, features 间的
交互都归因到 features 上。

ALE 很适合作为 Functional
decomposition.

exponential explosion
指数的爆炸。

practicability 实用性。

高阶的交互并不是很容易
可视化, 计算耗时会很让人
难受。

Functional ANOVA 要求 feature 之
间没有依赖

Generalized functional ANOVA
很难计算。

ALE 不提供方差分解。